

# 50.039 – Theory and Practice of Deep Learning

Alex

## Week 03: Pytorch basics I

[The following notes are compiled from various sources such as textbooks, lecture materials, Web resources and are shared for academic purposes only, intended for use by students registered for a specific course. In the interest of brevity, every source is not cited. The compiler of these notes gratefully acknowledges all such sources. ]

### Key content

- pytorch tensors: numpy with GPU transfer option
  - linear algebra similar to numpy
  - `torch.einsum` for general tensor multiplications with summing
  - data is stored in `.data` field algebra routine
- pytorch broadcasting rules
- **be able to write down mathematically what a certain pytorch operation does**
- **be able to decide what math formula can be realized with what pytorch linear**
- when one needs to use only data or handle gradients, tensor have `.data` and `.grad.data` fields

## 1 Pytorch tensor basics

[https://pytorch.org/tutorials/beginner/blitz/tensor\\_tutorial.html#sphx-glr-beginner-blitz-tensor-tutorial-py](https://pytorch.org/tutorials/beginner/blitz/tensor_tutorial.html#sphx-glr-beginner-blitz-tensor-tutorial-py)

Tensor mathematically:

- 1-tensor: a linear mapping  $v_1 \mapsto L(v_1)$ , representable as  $L(v_1) = u \cdot v_1$  by a vector  $u = (u_j)$
- 2-tensor: a bilinear mapping  $v_1, v_2 \mapsto L(v_1, v_2)$ , representable as  $L(v_1, v_2) = v_1^t A v_2 = \sum_{ij} v_{1,i} v_{2,j} A_{ij}$  by a matrix  $A = (A_{ij})$

- 3-tensor: a trilinear mapping  $v_1, v_2, v_3 \mapsto L(v_1, v_2, v_3)$ , representable as  $L(v_1, v_2, v_3) = \sum_{ijk} v_{1,i} v_{2,j} v_{3,k} A_{ijk}$  by a 3-dim array  $A = (A_{ijk})$
- n-tensor ... n-linear mapping ... representable by a n-dim array  $A = (A_{i_1 \dots i_n})$
- n-tensors  $\leftrightarrow$  n-dim arrays

Tensor in pytorch:

a representation of an numpy-array-like structure  $A_i$  or  $A_{i,j,k}$  or  $A_{i,j,k}$  or  $A_{i,j,k,l}$  with possibly more than 2 indices with benefits (for storing computed gradients).

## 1.1 init tensor as zeros, ones, constants

```
x= torch.empty((2,3)) %empty tensor
```

A tensor has three important properties:

- its `.size()` or `.shape`
- the dtype: its numerical type (most nns use `torch.float32`)
- device it is placed on (cpu, cuda:0, cuda:1)

getting its size: output is a `torch.Size()` object.

```
print(x.size())
print(x.shape)
```

Use list or tuple to get a list/tuple from that.

```
xs=tuple(x.size())
print(type(xs))
print(xs)
print(xs[0])
```

get its dtype:

```
print(x.dtype)
```

get its device placement

```
print(x.device)
```

`x.device` is a `torch.device` class

if you need strings, use `__repr__()`.

Test for equality with

```
x.device==torch.device('cuda:0')
x.dtype==torch.float %rhs is a torch.dtype object
x.dtype.__repr__()== 'torch.float32'
```

Important: you can print these anywhere in your execution code. no ugly fixed graph surprises.

```
x= torch.zeros((5,1))
y= torch.ones((5))
z= torch.empty((3,2,3))
a= a.new_full((3,2),42.) # tensor with a value, with same type and placement as tensor a
```

<https://pytorch.org/docs/stable/tensors.html> – dtypes and tensor types.  
dtype is the type of one element. tensor type depends on dtype and device placement

## 1.2 tensor from numpy

```
a=np.random.normal(5,size=(2,3)).astype('float32')
x=torch.tensor( a ) % this copies data
x2=torch.as_tensor(a) % this does NOT COPY data , and does nothing if its a tensor with ri
x3=torch.from_numpy( a ) % this does NOT COPY data - when this can be inappropriate? not re
```

## 1.3 tensor to numpy

```
t=x.data.numpy() #x.numpy() works only if x has no gradient attached )later(
```

## 1.4 change shape: reshape a tensor

```
a.view(...)

x=torch.ones((10))
x=x.view((-1,5))
```

Be careful which element lands where after reshaping!

## 1.5 change device: move tensor to gpu / cpu device

```
device=torch.device('cuda:0')
xg=x.to(device)
xc=x.to(torch.device('cpu'))
```

## 1.6 change tensor dtype

<https://pytorch.org/docs/stable/tensors.html>

print your type

```
print(x.type())
```

cast tensor

```
x=x.type(torch.FloatTensor) % pretty common error source in dataloaders
x=x.type_as(a) # a is another tensor
```

## 1.7 create tensor of same type and device as another tensor

This is useful when one uses `torch.nn.DataParallel` for multi-GPU computations – then one does not want to instantiate a tensor explicitly on a fixed GPU like `cuda:0`.

```
b=a.new_ones((3,2))
```

### Debugging in pytorch

Most of my own programming errors come from above three properties: try to compute results

- with incompatible shapes.
- from two tensors with incompatible numerical type (integer and float, float32 and double),
- with incompatible devices (one on cpu, other on GPU),

The good news: in pytorch you can print `.size()` anywhere in the running code, also its dtype and its device placement

debugging advice: RTFM and print the shapes.

## 2 broadcasting

<https://pytorch.org/docs/stable/notes/broadcasting.html>

```
a = torch.ones((4))
b = torch.ones((1,4))
torch.add(a,b) → (1,4)
```

```
a = torch.ones((4))
b = torch.ones((4,1))
torch.add(a,b) → (4,4)!!!
```

```
a = torch.ones((3))
b = torch.ones((4,1))
torch.add(a,b) → (4,3)
```

```

a = torch.ones((3))
b = torch.ones((1,4))
torch.add(a,b) → ERR

```

- smaller tensor gets filled **from the left** with singleton dimensions until he has same dimensionality as larger tensor, as if `.unsqueeze(0)` would be applied again and again

```

a = torch.ones((2,5))
b = torch.ones((5))
torch.add(a,b) → (2,5)
→ a = torch.ones((2,5)) ok as is
→ b = torch.ones((1,5)) copy 1x until (2,5)

```

```

a = torch.ones((3))
b = torch.ones((4,1))
torch.add(a,b) → (4,3)
→ a = torch.ones((1,3)) copy 3x until (4,3)
→ b = torch.ones((4,1)) copy 2x until (4,3)

```

- whenever a dimension with size 1 meets a dimension with a size  $k > 1$ , then the smaller vector is replicated/copied  $k - 1$  times in this dimension until he reaches in this dimension size  $k$

## 2.1 linear algebra: sum, inner product, matrix product

`torch.mm(a,b)` dot product, not broadcasting.  $a, b$  must be 1-tensors

$$\begin{aligned}
 a.size() &= (d) \\
 b.size() &= (d) \\
 torch.dot(a,b) &= \sum_{d'} a_{d'} b_{d'} = \sum_{d'} a[d'] b[d'] \rightarrow torch.dot(a,b).size() = ()
 \end{aligned}$$

`torch.mm(A,B)` matrix multiplication, not broadcasting.  $A, B$  must be 2-tensors

$$\begin{aligned}
A.size() &= (i, k) \\
B.size() &= (k, l) \\
torch.mm(A, B)[i, l] &= \sum_{k'} A_{i, k'} B_{k', l} = \sum_k A[i, k'] B[k', l] \rightarrow torch.mm(A, B).size() = (i, l)
\end{aligned}$$

`torch.bmm(A, B)` **batched** matrix multiplication, not broadcasting.  $A, B$  must be 3-tensors. multiplication along last dim of  $A$  and second dim of  $B$ .

$$\begin{aligned}
A.size() &= (b, i, k) \\
B.size() &= (b, k, l) \\
torch.bmm(A, B)[b, i, l] &= \sum_{k'} A_{b, i, k'} B_{b, k', l} = \sum_k A[b, i, k'] B[b, k', l] \rightarrow torch.bmm(A, B).size() = (b, i, l)
\end{aligned}$$

`torch.bmm(A, B)` performs for every index  $k$  a matrix multiplication between  $A[k, :, :]$  and  $B[k, :, :]$   
– its a for loop over  $k$  of `torch.mm(A[k, :, :], B[k, :, :])`

Think: `torch.bmm(A, B)` given a known shape of  $A$  puts what restrictions on  $B$ ??

## 2.2 linear algebra: shapes dont fit?!

`torch.squeeze(A, dim=2)` - remove singleton dim  $(a, b, 1, c) \rightarrow (a, b, c)$   
`torch.unsqueeze(A, dim=1)` - insert singleton dim  $(a, b, c) \rightarrow (a, 1, b, c)$   
`torch.unsqueeze(A, dim=0)` - insert singleton dim  $(a, b, c) \rightarrow (1, a, b, c)$

example: want to compute with mm matrix vector product  $(vA)_l = \sum_k v_k A_{k, l}$ .  $v$  is 1-tensor, so cannot use `torch.mm(v, A)`. So add a singleton dimension in  $v$ :

$$\begin{aligned}
torch.mm(v.unsqueeze(0), A) &\rightarrow (1, L) \\
torch.mm(v.unsqueeze(0), A).squeeze(0) &\rightarrow (1, L) \rightarrow (L)
\end{aligned}$$

`torch.transpose(A, dim1, dim2)` swaps two dimensions  
`torch.Tensor.permute(*dims)` permutes a set of dimensions rather than just swapping two

## 2.3 matmul the monster

`torch.matmul(A, B)` - matrix multiplication with broadcasting - that is only 1 dimension is summed out <https://pytorch.org/docs/stable/torch.html#torch.matmul> This function is a bit tricky, performs broadcasting of shapes (<https://pytorch.org/docs/stable/notes/broadcasting.html>) - for tensors which are  $N \geq 3$ -tensors broadcasting is done on all dimensions except

on the last two dimensions, then multiplies along the last dimension of  $A$ , and over the second last dimension of  $B$  – `torch.einsum` is more clear here and helps if `matmul` is not clear.

The non-matrix (i.e. batch) dimensions are broadcasted (and thus must be broadcastable).

#### Broadcasting warning

You cannot avoid getting to know the broadcasting rules (quiz?).

many simple functions like `+`, `*` do broadcasting by default.

`torch.tensordot(A,B, dims=(list1,list2) )` - general tensor contractions with broadcasting – that is multiple dimensions can be summed out <https://pytorch.org/docs/stable/torch.html#torch.tensordot>

## 2.4 einsum is the monster ... ?

#### `torch.einsum`

a general way to do all kinds of batched and non-batched tensor multiplications: `torch.einsum`

<https://rockt.github.io/2018/04/30/einsum>

rule:

left of `– >`: all tensors separated by `,` which are to be multiplied and summed.

**indices that have same name in multiple tensors, will get multiplied together**

right of `– >` the result tensor with remaining indices. **All indices missing right of `– >` are summed out so that they vanish in the result.**

## 2.5 other useful stuff

`res=torch.where(x>5,x,y)`

<https://pytorch.org/docs/stable/tensors.html>